

CME 193: Introduction to Scientific Python

Lecture 6: Numpy, Scipy, Matplotlib

Nolan Skochdopole

`stanford.edu/class/cme193`

Contents

- Homeworks and Project
- Numpy
- Scipy
- Matplotlib
- Exercises

Homeworks

Homework 1 was due last night. Hope that went well.

I will post homework 2 tomorrow - will be longer than the first.

Will be due roughly two weeks later - will post exact date with it.

Project

Welcome to do your own project instead of the homework.

If you do, please give a brief proposal (1-2 paragraphs) by this Friday.

Project will be due same time as homework 2.

Contents

- Homeworks and Project
- **Numpy**
- Scipy
- Matplotlib
- Exercises

Numpy

- Fundamental package for scientific computing with Python
- N-dimensional array object
- Linear algebra, Fourier transform, random number capabilities
- Building block for other packages (e.g. Scipy)
- Open source

Numpy

- Fundamental package for scientific computing with Python
- N-dimensional array object
- Linear algebra, Fourier transform, random number capabilities
- Building block for other packages (e.g. Scipy)
- Open source

import numpy as np

Basics:

```
import numpy as np

A = np.array([[1, 2, 3], [4, 5, 6]])
print A
# [[1 2 3]
#  [4 5 6]]
print A[0,0] # 1
print A[0,1:3] # [2 3]

Af = np.array([1, 2, 3], float)
```

Slicing as usual.

More basics

```
np.arange(0, 1, 0.2)
# array([ 0. ,  0.2,  0.4,  0.6,  0.8])

np.linspace(0, 2*np.pi, 4)
# array([ 0.0,  2.09,  4.18,  6.28])

A = np.zeros((2,3))
# array([[ 0.,  0.,  0.],
#        [ 0.,  0.,  0.]])
# np.ones, np.diag
A.shape
# (2, 3)
```

More basics

```
np.random.random((2,3))
# array([[ 0.78084261,  0.64328818,  0.55380341],
#        [ 0.24611092,  0.37011213,  0.83313416]])

a = np.random.normal(loc=1.0, scale=2.0, size=(2,2))
# array([[ 2.87799514,  0.6284259 ],
#        [ 3.10683164,  2.05324587]])

np.savetxt("a_out.txt", a)
# save to file
b = np.loadtxt("a_out.txt")
# read from file
```

Arrays are mutable

```
A = np.zeros((2, 2))  
# array([[ 0.,  0.],  
#        [ 0.,  0.]])  
C = A  
C[0, 0] = 1  
  
print A  
# [[ 1.  0.]  
#   [ 0.  0.]])
```

Array attributes

```
a = np.arange(10).reshape((2,5))  
  
a.ndim      # 2 dimension  
a.shape     # (2, 5) shape of array  
a.size      # 10 # of elements  
a.T         # transpose  
a.dtype     # data type
```

Basic operations

Arithmetic operators: **elementwise** application

```
a = np.arange(4)
# array([0, 1, 2, 3])

b = np.array([2, 3, 2, 4])

a * b # array([ 0,  3,  4, 12])
b - a # array([2, 2, 0, 1])

c = [2, 3, 4, 5]
a * c # array([ 0,  3,  8, 15])
```


Also, we can use += and *+=.

Array broadcasting

When operating on two arrays, numpy compares shapes. Two dimensions are compatible when

1. They are of equal size
2. One of them is 1

Array broadcasting



Array broadcasting with scalars

This also allows us to add a constant to a matrix or multiply a matrix by a constant

```
A = np.ones((3,3))  
  
print 3 * A - 1  
# [[ 2.  2.  2.]  
#   [ 2.  2.  2.]  
#   [ 2.  2.  2.]
```


Vector operations

- inner product
- outer product
- dot product (matrix multiplication)

```
# note: numpy automatically converts lists
u = [1, 2, 3]
v = [1, 1, 1]

np.inner(u, v)
# 6
np.outer(u, v)
# array([[1, 1, 1],
#        [2, 2, 2],
#        [3, 3, 3]])
np.dot(u, v)
# 6
```

Matrix operations

First, define some matrices:

```
A = np.ones((3, 2))
# array([[ 1.,  1.],
#        [ 1.,  1.],
#        [ 1.,  1.]])
A.T
# array([[ 1.,  1.,  1.],
#        [ 1.,  1.,  1.]])

B = np.ones((2, 3))
# array([[ 1.,  1.,  1.],
#        [ 1.,  1.,  1.]])
```

Matrix operations

```
np.dot(A, B)
# array([[ 2.,  2.,  2.],
#        [ 2.,  2.,  2.],
#        [ 2.,  2.,  2.]])

np.dot(B, A)
# array([[ 3.,  3.],
#        [ 3.,  3.]])

np.dot(B.T, A.T)
# array([[ 2.,  2.,  2.],
#        [ 2.,  2.,  2.],
#        [ 2.,  2.,  2.]])

np.dot(A, B.T)
# Traceback (most recent call last):
#   File "<stdin>", line 1, in <module>
# ValueError: shapes (3,2) and (3,2) not aligned: ...
# ... 2 (dim 1) != 3 (dim 0)
```

Operations along axes

```
a = np.random.random((2,3))
# array([[ 0.9190687 ,  0.36497813,  0.75644216],
#        [ 0.91938241,  0.08599547,  0.49544003]])
a.sum()
# 3.5413068994445549
a.sum(axis=0) # column sum
# array([ 1.83845111,  0.4509736 ,  1.25188219])
a.cumsum()
# array([ 0.9190687 ,  1.28404683,  2.04048899,  2.9598714 ,
#        3.04586687,  3.5413069 ])
a.cumsum(axis=1) # cumulative row sum
# array([[ 0.9190687 ,  1.28404683,  2.04048899],
#        [ 0.91938241,  1.00537788,  1.50081791]])
a.min()
# 0.0859954690403677
a.max(axis=0)
# array([ 0.91938241,  0.36497813,  0.75644216])
```

Slicing arrays

More advanced slicing

```
a = np.random.random((4,5))  
  
a[2, :]  
# third row, all columns  
a[1:3]  
# 2nd, 3rd row, all columns  
a[:, 2:4]  
# all rows, columns 3 and 4
```

Iterating over arrays

- Iterating over multidimensional arrays is done with respect to the first axis: `for row in A`
- Looping over all elements: `for element in A.flat`

Reshaping

Reshape using `reshape`. Total size must remain the same.

Resize using `resize`, always works: chopping or appending zeros

First dimension has 'priority', so beware of unexpected results

Try it!

Reshaping

Reshape using `reshape`. Total size must remain the same.

Resize using `resize`, always works: chopping or appending zeros

First dimension has 'priority', so beware of unexpected results

Try it!

Reshaping

Reshape using `reshape`. Total size must remain the same.

Resize using `resize`, always works: chopping or appending zeros

First dimension has 'priority', so beware of unexpected results

Try it!

Matrix operations

<code>eye(3)</code>	Identity matrix
<code>trace(A)</code>	Trace
<code>column_stack((A,B))</code>	Stack column wise
<code>row_stack((A,B,A))</code>	Stack row wise

Linear algebra

```
import numpy.linalg
```

<code>qr</code>	Computes the QR decomposition
<code>cholesky</code>	Computes the Cholesky decomposition
<code>inv(A)</code>	Inverse
<code>solve(A,b)</code>	Solves $Ax = b$ for A full rank
<code>lstsq(A,b)</code>	Solves $\arg \min_x \ Ax - b\ _2$
<code>eig(A)</code>	Eigenvalue decomposition
<code>eig(A)</code>	Eigenvalue decomposition for symmetric or hermitian
<code>eigvals(A)</code>	Computes eigenvalues.
<code>svd(A, full)</code>	Singular value decomposition
<code>pinv(A)</code>	Computes pseudo-inverse of A

Fourier transform

```
import numpy.fft
```

- `fft` 1-dimensional DFT
- `fft2` 2-dimensional DFT
- `fftn` N-dimensional DFT
- `ifft` 1-dimensional inverse DFT (etc.)
- `rfft` Real DFT (1-dim)
- `ifft` Imaginary DFT (1-dim)

Random sampling

```
import numpy.random
```

<code>rand(d0,d1,...,dn)</code>	Random values in a given shape
<code>randn(d0, d1, ...,dn)</code>	Random standard normal
<code>randint(lo, hi, size)</code>	Random integers [lo, hi)
<code>choice(a, size, repl, p)</code>	Sample from a
<code>shuffle(a)</code>	Permutation (in-place)
<code>permutation(a)</code>	Permutation (new array)

Distributions in random

```
import numpy.random
```

The list of distributions to sample from is quite long, and includes

- beta
- binomial
- chisquare
- exponential
- dirichlet
- gamma
- laplace
- lognormal
- pareto
- poisson
- power

Exercise

Write a script that creates a random square matrix (A) with standard normal random variables and a random column vector (b) of the same size (also standard normal random variables).

Solve the system $Ax = b$.

Compute both the 2-norm and infinity norm of x as well as the Frobenius norm of A .

Hint see documentation for `np.linalg.norm`

Exercise

```
import numpy as np

n = 200
A = np.random.randn(n,n)
b = np.random.randn(n,1)
x = np.linalg.solve(A,b)

print np.linalg.norm(x)
print np.linalg.norm(x, np.inf)
print np.linalg.norm(A, 'fro')
```


Contents

- Homeworks and Project
- Numpy
- **Scipy**
- Matplotlib
- Exercises

What is SciPy?

SciPy is a library of algorithms and mathematical tools built to work with NumPy arrays.

- linear algebra - *scipy.linalg*
- statistics - *scipy.stats*
- optimization - *scipy.optimize*
- sparse matrices - *scipy.sparse*
- signal processing - *scipy.signal*
- etc.

Scipy Linear Algebra

Slightly different from `numpy.linalg`. Always uses BLAS/LAPACK support, so could be faster.

Some more functions.

Functions can be slightly different.

Scipy Optimization

- General purpose minimization: CG, BFGS, least-squares
- Constrained minimization; non-negative least-squares
- Minimize using simulated annealing
- Scalar function minimization
- Root finding
- Check gradient function
- Line search

Scipy Statistics

- Mean, median, mode, variance, kurtosis
- Pearson correlation coefficient
- Hypothesis tests (ttest, Wilcoxon signed-rank test, Kolmogorov-Smirnov)
- Gaussian kernel density estimation

See also SciKits (or scikit-learn).

Scipy sparse

- Sparse matrix classes: CSC, CSR, etc.
- Functions to build sparse matrices
- `sparse.linalg` module for sparse linear algebra
- `sparse.csgraph` for sparse graph routines

Scipy signal

- Convolutions
- B-splines
- Filtering
- Continuous-time linear system
- Wavelets
- Peak finding

Scipy IO

Methods for loading and saving data

- Matlab files
- Matrix Market files (sparse matrices)
- Wav files

Example

```
from scipy import optimize

def f(x):
    return [x[0] + 0.5 * (x[0] - x[1])**3 - 1.0,
            0.5 * (x[1] - x[0])**3 + x[1]]

x0 = [0, 0] # initial guess

sol = optimize.root(f, x0)

print sol.x
print sol.success
```

Exercise

Create a matrix (A) of random entries (your choice on distribution) with $m > n$ (more rows than columns).

Create a column vector $b \in \mathbb{R}^m$.

Find x that minimizes $\|Ax - b\|_2$. What is the norm of the residual?

Hint: use `scipy.linalg.lstsq`

Exercise

```
import numpy as np
from scipy import linalg

n = 100
m = 200

A = np.random.randn(m,n)
b = np.random.randn(m,1)

x = linalg.lstsq(A, b)
print linalg.norm(np.dot(A,x[0])-b)
```

Contents

- Homeworks and Project
- Numpy
- Scipy
- **Matplotlib**
- Exercises


What is Matplotlib?

- Plotting library for Python
- Works well with Numpy
- Syntax similar to Matlab

Scatter Plot

```
import numpy as np
import matplotlib.pyplot as plt


x = np.linspace(0, 10, 1000)
y = np.power(x, 2)
plt.plot(x, y)
plt.show()
```



Seaborn makes plot pretty

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

x = np.linspace(0, 10, 1000)
y = np.power(x, 2)
plt.plot(x, y)
plt.show()
```



`plt.show()`

Calling `plt.show()` displays the figure objects to screen.

To create/display multiple figures, use `plt.figure()`, very similar to Matlab.

plt.figure()

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

x = np.linspace(0, 10, 1000)
y1 = np.power(x, 2)
y2 = np.power(x, 3)
f1 = plt.figure()
plt.plot(x, y1)
f2 = plt.figure()
plt.plot(x, y2)
plt.show()
```

If we had not used separate figure calls, it would default to plotting both curves on same figure.

Scatter Plot

Adding titles and labels


```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

f, ax = plt.subplots(1, 1, figsize=(5,4))

x = np.linspace(0, 10, 1000)
y = np.power(x, 2)
ax.plot(x, y)
ax.set_xlim((1, 5))
ax.set_ylim((0, 30))
ax.set_xlabel('my x label')
ax.set_ylabel('my y label')
ax.set_title('plot title, including  $\Omega$ ')

plt.tight_layout()
plt.savefig('line_plot_plus.pdf')
```

Scatter Plot



Scatter Plot


Adding multiple lines and a legend

```
x = np.linspace(0, 10, 50)
y1 = np.power(x, 2)
y2 = np.power(x, 3)

plt.plot(x, y1, 'b-', label='$x^2$')
plt.plot(x, y2, 'go', label='$x^3$')
plt.xlim((1, 5))
plt.ylim((0, 30))
plt.xlabel('my x label')
plt.ylabel('my y label')
plt.title('plot title, including $\Omega$')
plt.legend()

plt.savefig('line_plot_plus2.pdf')
```

Scatter Plot



Histogram

```
data = np.random.randn(1000)


f, (ax1, ax2) = plt.subplots(1, 2, figsize=(6,3))

# histogram (pdf)
ax1.hist(data, bins=30, normed=True, color='b')

# empirical cdf
ax2.hist(data, bins=30, normed=True, color='r',
          cumulative=True)

plt.savefig('histogram.pdf')
```

Histogram



Box Plot

```
samp1 = np.random.normal(loc=0., scale=1., size=100)
samp2 = np.random.normal(loc=1., scale=2., size=100)
samp3 = np.random.normal(loc=0.3, scale=1.2, size=100)

f, ax = plt.subplots(1, 1, figsize=(5,4))

ax.boxplot((samp1, samp2, samp3))
ax.set_xticklabels(['sample 1', 'sample 2', 'sample 3'])
plt.savefig('boxplot.pdf')
```


Box Plot





Image Plot

```
A = np.random.random((100, 100))  
  
plt.imshow(A)  
plt.hot()  
plt.colorbar()  
  
plt.savefig('imageplot.pdf')
```

Image Plot




Wire Plot

matplotlib toolkits extend functionality for other kinds of visualization


```
from mpl_toolkits.mplot3d import axes3d  
  
ax = plt.subplot(111, projection='3d')  
X, Y, Z = axes3d.get_test_data(0.1)  
ax.plot_wireframe(X, Y, Z, linewidth=0.1)  
  
plt.savefig('wire.pdf')
```

Wire Plot



Possibilities

A lot is possible, but not always easy to figure out how...



Exercise

Plot the following function on the interval $[-2, 2]$.

$$f(x) = \sin(x) \cos(x - \pi) e^{-x}$$

Exercise

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

def f(x):
    return np.sin(x)*np.cos(x-np.pi)*np.exp(-1.*x)

x = np.linspace(-2,2)
y = f(x)
plt.plot(x, y)
plt.show()
```


Contents

- Homeworks and Project
- Numpy
- Scipy
- Matplotlib
- Exercises

Exercises

See course website for exercises for this lecture.